

# Key-Value Store Databases

Advanced Databases  
INFO - H 415

**Presented by:**  
**Mayr Dionisius**  
**Elezi Herma**  
**Suau Thomas**  
**İşlek Rana**

# Table of Content

- 01 Introduction
- 02 Key-Value Store
- 03 Redis
- 04 etcd
- 05 Methodology
- 06 Application implementation
- 07 Conclusion

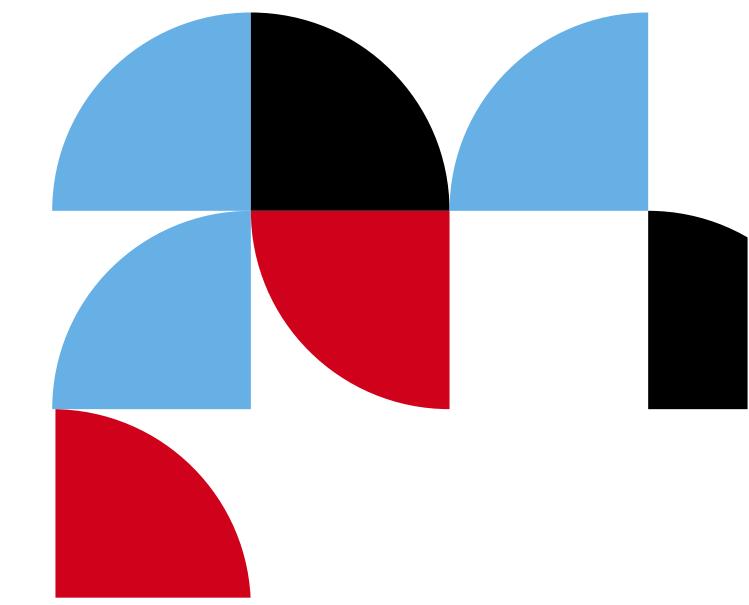
**~ Part 1 ~**

# **Introduction to technology:**

## **Key-Value Storing**

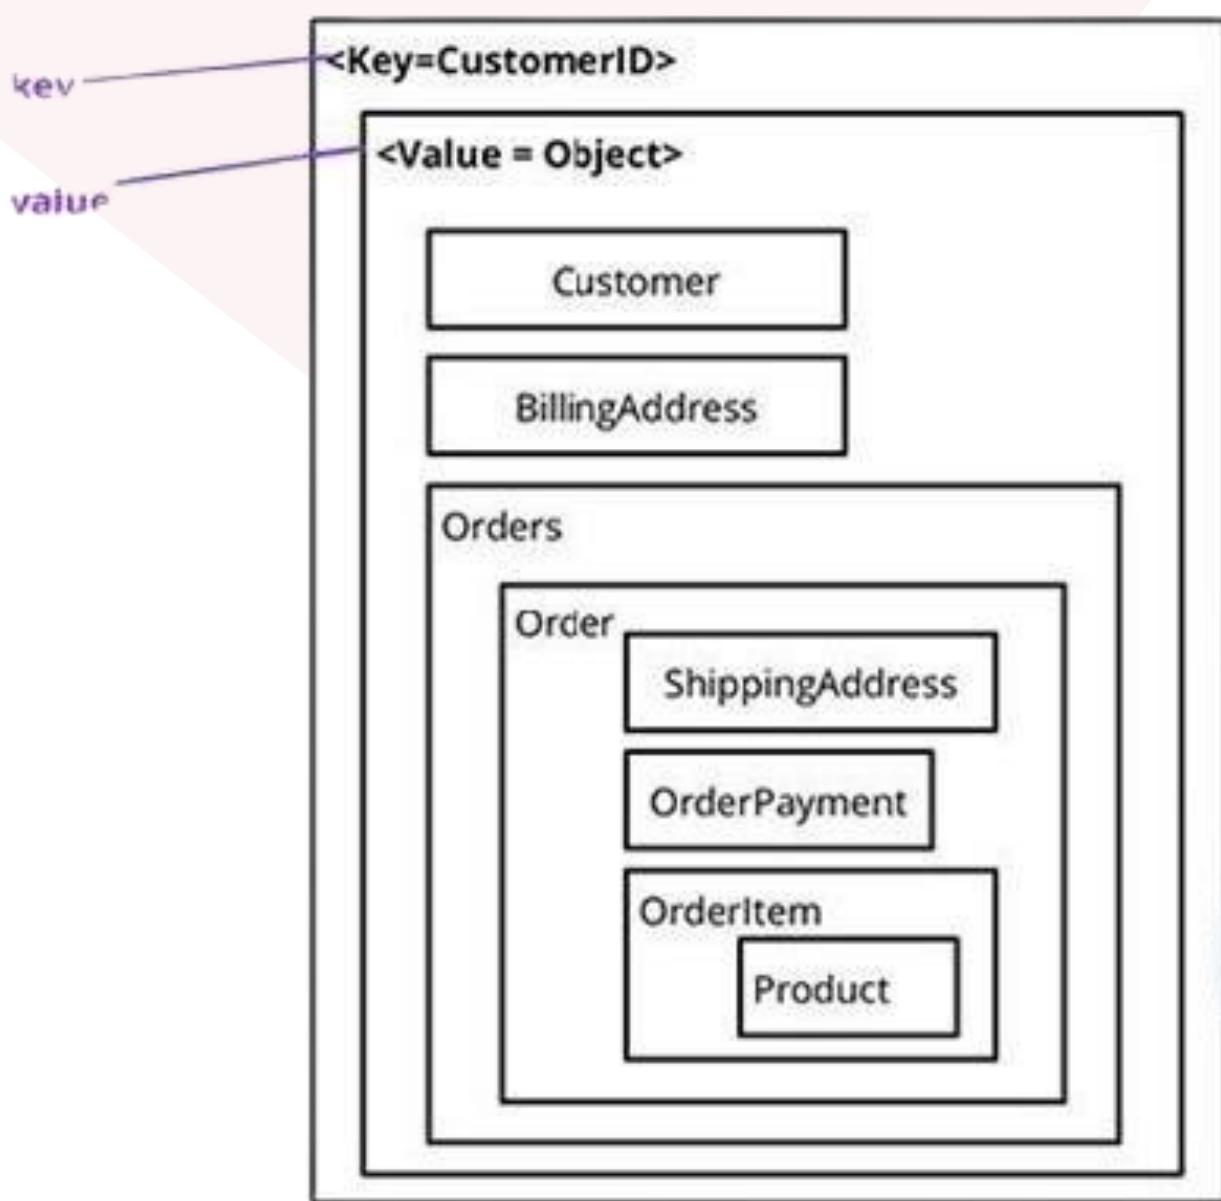
# Introduction

- Our project aims to select the right Database Management System for a caching application;
- Two technologies are being compared:
  - Traditional relational databases
  - Key-value stores
- We aim to answer which technology would be more suitable for our application.



# Key-Value Stores

Key-value stores pair unique keys with specific values. They're schema-less for data flexibility, and are known for fast data retrieval, scalability, and simple data models.



# Key-Value Stores

## Importance in Modern Database Systems

- Crucial Role
- Large Volumes of Data
- Quick and Efficient Data Access
- Use Cases
- Simplicity

# Key-Value Stores

## Advantages



High Performance



Scalability



Flexibility

# Key-Value Stores

## Use Cases



Caching



Session Storage

**~ Part 2 ~**

**Introduction to the two tools:  
Redis and etcd**

**+**

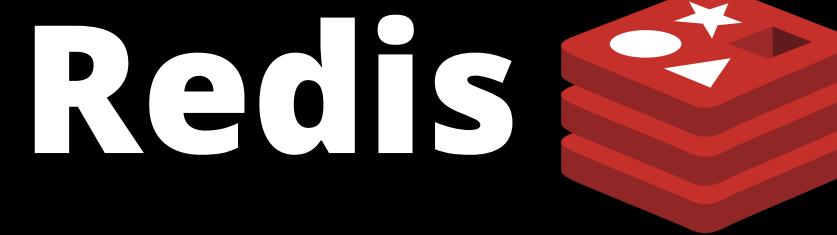
**Advantages and Disadvantages**

# Redis



## ReThink Caching

- **Remote Dictionary Server**
- **Open-source**
- **In-memory data structure store**
- **Various data structures:**
  - **strings**
  - **hashes**
  - **lists**
  - **sets**
  - **...**
- **Designed for high performance:**
  - **Low-latency access**



## FEATURES

### Data Structures

Redis supports a variety of data providing flexibility for modeling

### Persistence

Redis offers options for persistence, allowing data to be saved to disk for durability

### Replication

Redis supports master-slave replication, enabling redundancy and high availability.

### Partitioning

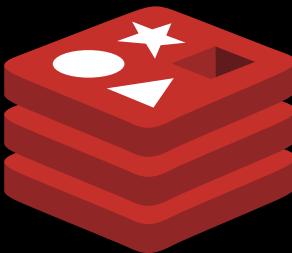
Redis can be horizontally scaled through sharding, distributing data across multiple nodes.



## (MOSTLY) A SINGLE THREADED APPLICATION

- Generally IO bound -> Parallelism wouldn't help in many cases
- It makes it simple
  - No thread synchronization (fast!)
  - Lockless atomic actions
  - Easier to maintain
- Where is it multi-thread?
  - Some small thread safe contexts are being moved to separate threads
  - Still, at any time Redis will still be serving a **single** request.
- Sharding can help CPU-bound applications

# Redis



hello world	String
011011010110111101101101	Bitmap
{23334}{6634728}{916}	Bitfield
{a: "hello", b: "world"}	Hash
[A>B>C>C]	List
{A<B<C}	Set
{A:1, B:2, C:3}	Sorted set
{A: (50.1, 0.5)}	Geospatial
01101101 01101111 01101101	Hyperlog
{id1=time1.seq({ a: "foo", a: "bar" }) }	Stream

Source:

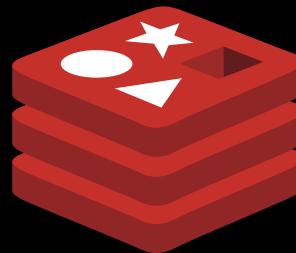
<https://architecturenotes.cc/redis/>



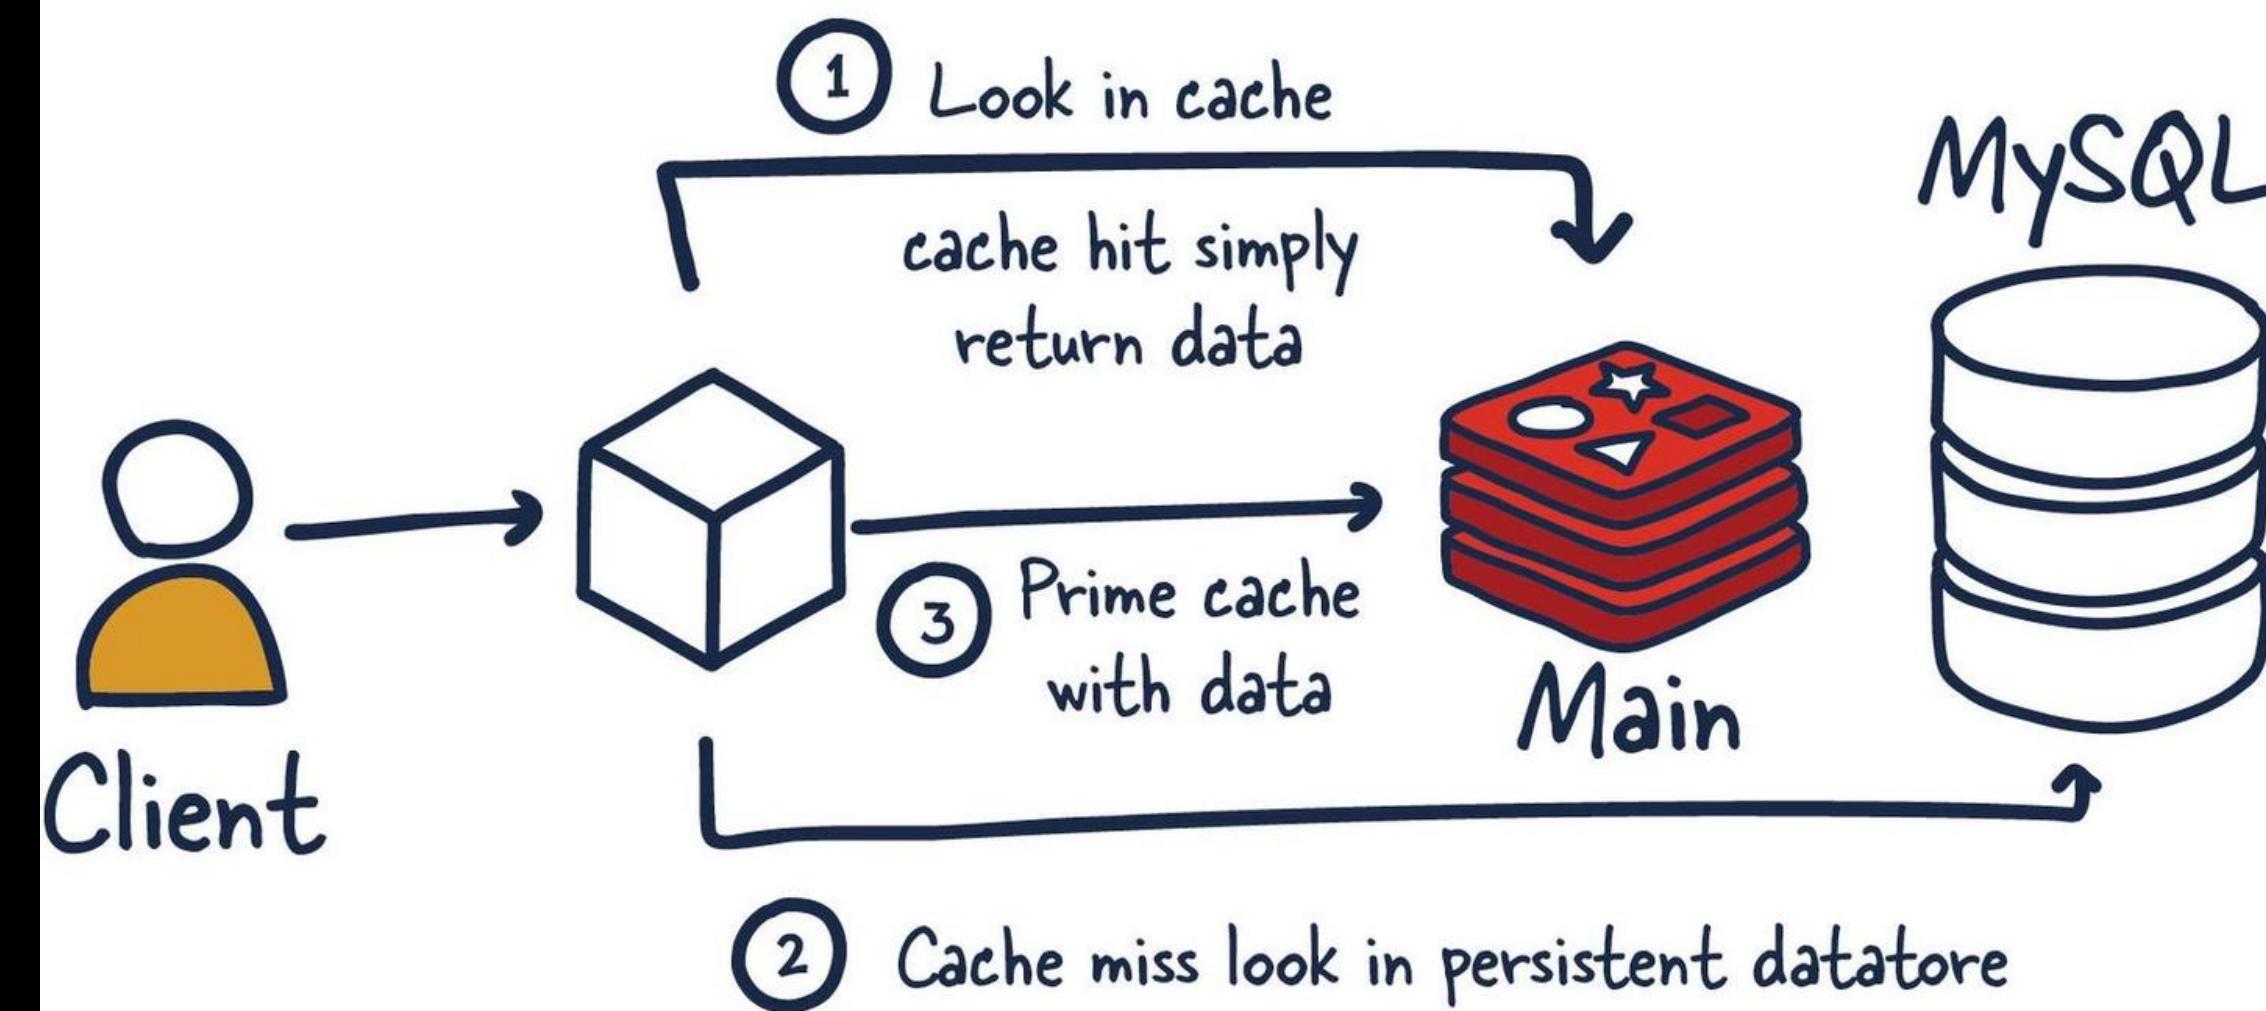
## SIMPLE

- Initial version developed in a short amount of time
- Easy to use
  - SET bike:1 "My Bike"
  - HSET bike:1 model Deimos brand Ergonom type 'Enduro bikes' price 4972
  - HGET bike:1 price
  - BITFIELD pla:1:stats INCRBY u32 #0 50 INCRBY u32 #11
  - EXPIRE bike:1 10
- Quite popular

# Redis



## How is redis traditionally used



Source:

<https://architecturenotes.co/redis/>

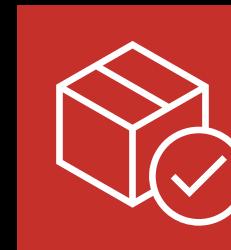


## USE CASES



### Caching

Due to its in-memory nature, Redis is widely used as a caching layer to store frequently accessed data for quick retrieval



### Session Storage

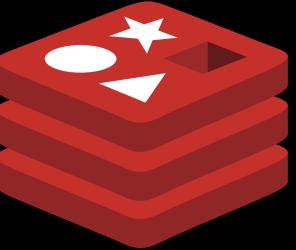
Storing session data in Redis ensures fast access and efficient management of user sessions in web applications



### Autocomplete

The fast read and write capabilities, together with the data structures provided make Redis suitable for auto-completing applications (such as recent contacts)

# Redis



## STRENGTHS

### Performance

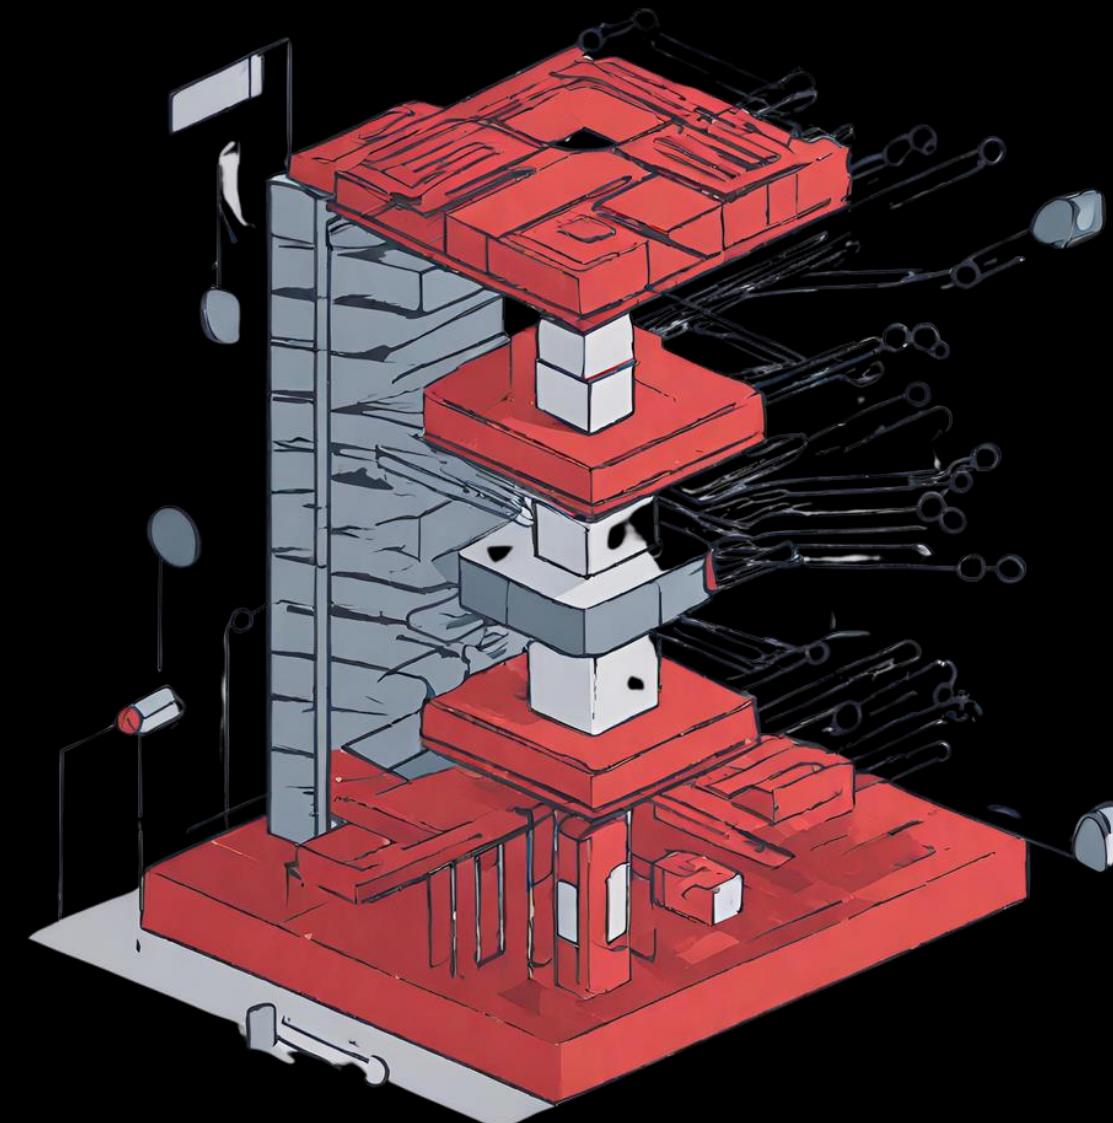


Redis excels in terms of performance, with low-latency access to data.

### Versatility

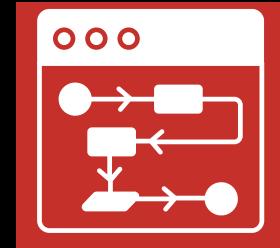


Leveraging the data structures provided makes it versatile enough to be used in multiple use cases.



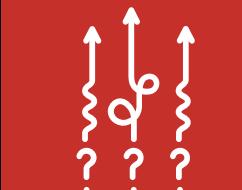
## WEAKNESSES

### Single-threaded



For some workloads, the single-threaded nature may become a bottleneck.

### Complex Queries



Redis may not be the best choice for scenarios that involve complex queries, as it prioritizes simplicity and speed.



# etcd

## Distributed Reliable Key-Value Store

A distributed, reliable key-value store for the most critical data of a distributed system



## WHAT IS etcd?

etcd, short for "**distributed et cetera**," is an open-source distributed key-value store designed for reliability, consistency, and simplicity in distributed systems.

The distributed key-value store etcd is specifically designed for managing configuration data, coordinating distributed systems (clusters), and ensuring strong consistency through **the Raft consensus algorithm**.

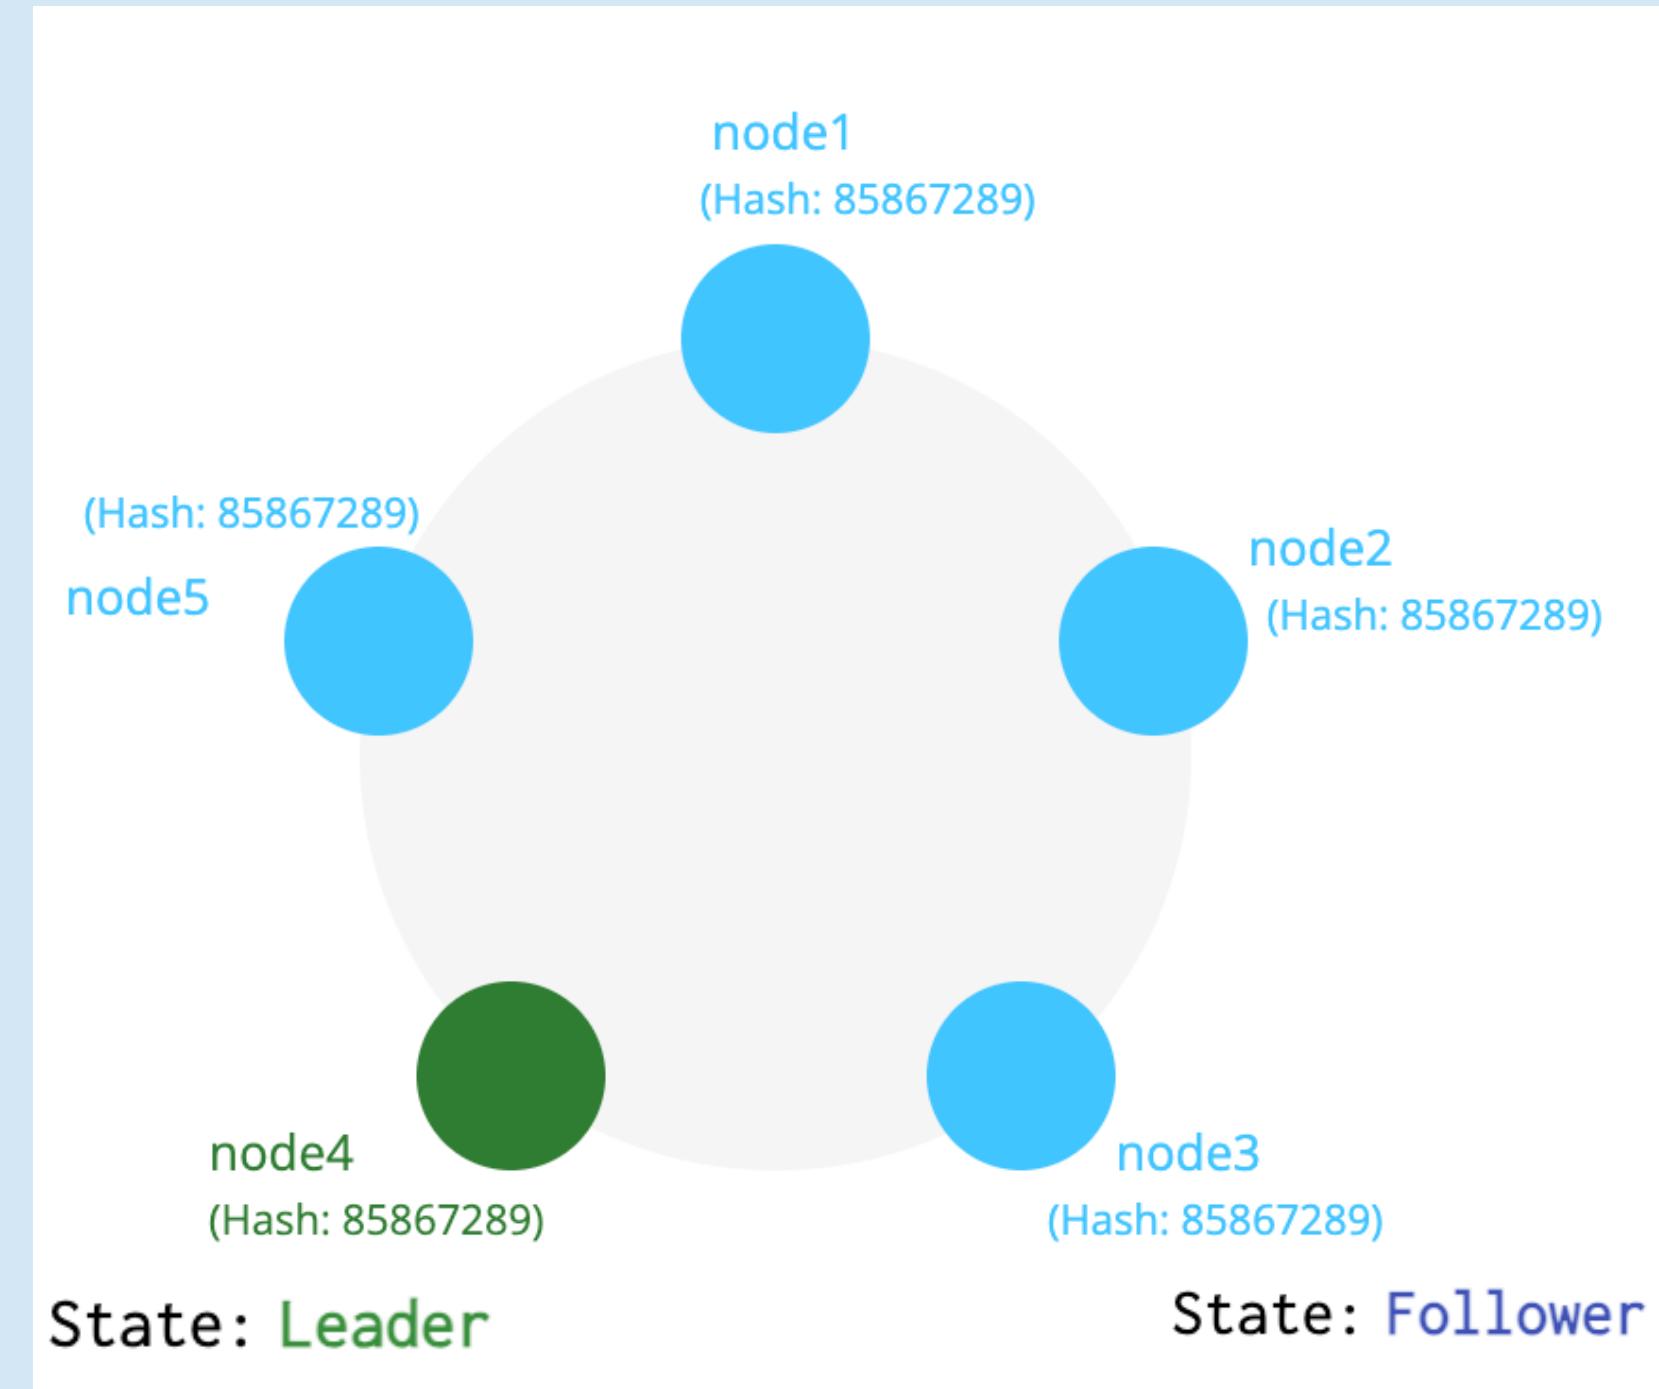
It's the core component of *Kubernetes* -> mainly used stores and manage state data, configuration data and metadata.

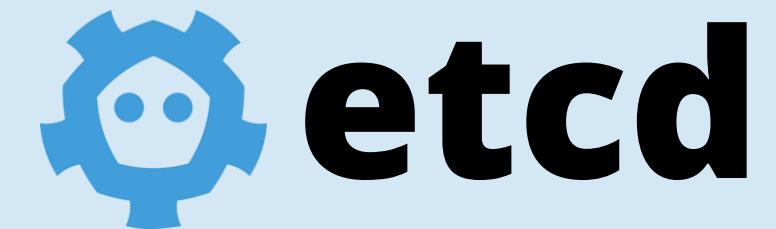


# Architecture

- etcd is built around the **Raft consensus algorithm** -> providing a fault-tolerant and consistent distributed key-value store.
- It operates in a cluster model to maintain data consistency.
- The Raft algorithm ensures that a **leader** is elected to coordinate and *replicate* data across the cluster

The architecture is **very robust**, but there is a *loss in IO write speed*.





## Key Features

### Raft Consensus Algorithm

etcd employs the Raft algorithm to ensure strong consistency and fault tolerance, providing a reliable foundation for distributed systems.

### Distributed Configuration

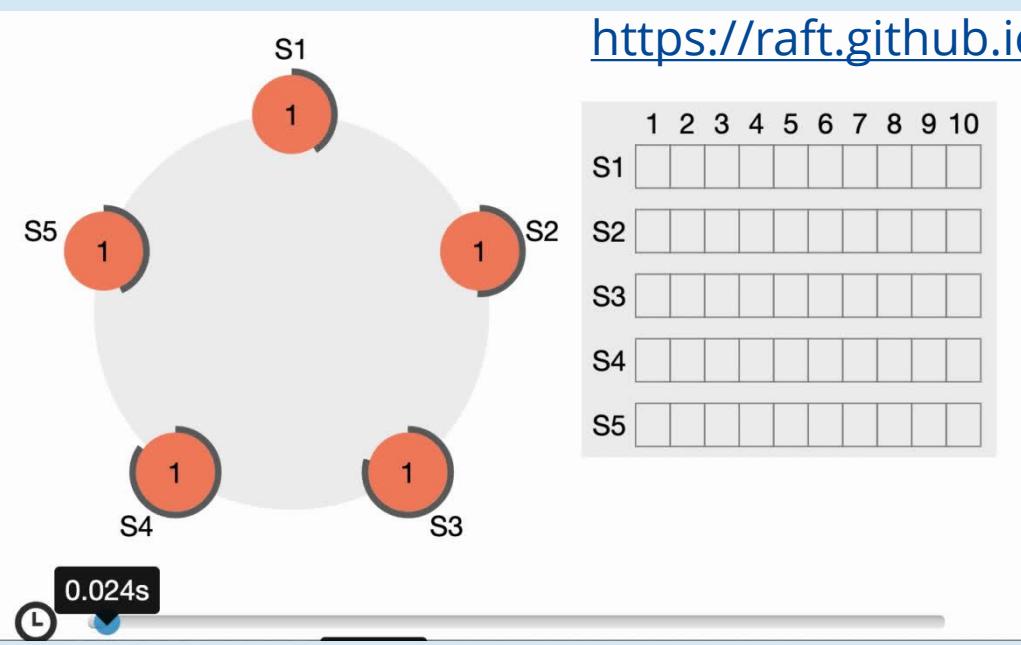
etcd excels in managing dynamic configurations across distributed applications, ensuring that all nodes have access to the latest configuration settings.

### Atomic Transactions

etcd supports atomic transactions, allowing a series of operations to be executed as a single, indivisible unit based on If/Then/Else block statements.

### Watch Notifications

Developers can watch for changes in etcd key-value pairs, enabling real-time notifications when configuration data is updated.



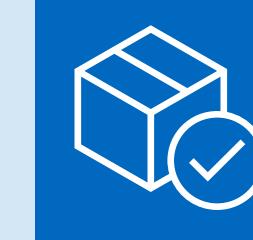


## Use Cases



### Kubernetes clusters management

etcd is the core component of Kubernetes to store and manage state data, configuration data and metadata.



### Service Discovery and Load Balancing

etcd can be used as a central registry to store information about available services, their locations, and health status. Each service instance registers itself in etcd, making it discoverable by other services.



### Distributed Locking and Coordination

etcd can coordinate distributed systems by providing a consistent view of shared data and managing distributed locks.

## Strengths and Weaknesses

### Consistency and reliability



etcd ensures strong consistency through the Raft consensus algorithm, making it reliable for critical data management.

### Fault Tolerance



distributed nature of etcd, combined with Raft, provides fault tolerance, ensuring system reliability in the face of node failures.



### Performance for Key-Value Storage



While etcd is robust in consistency and coordination, it may not match the raw performance of in-memory key-value stores like Redis, especially in scenarios prioritizing low-latency access.

### Specialized Use Case



etcd is designed for specific use cases. etcd won't be efficient for other application purposes.

**~ Part 3 ~**

# **Methodology and Sample Application**



# Methodology

# Methodology

## Yahoo! Cloud Serving Benchmark (YCSB)

- Industry standard benchmark for over a decade
  - Focus on key-value and cloud serving stores
- Simulates multiple load scenarios
  - Time and throughput
- go-ycsb wrapper



# Methodology

## Workloads

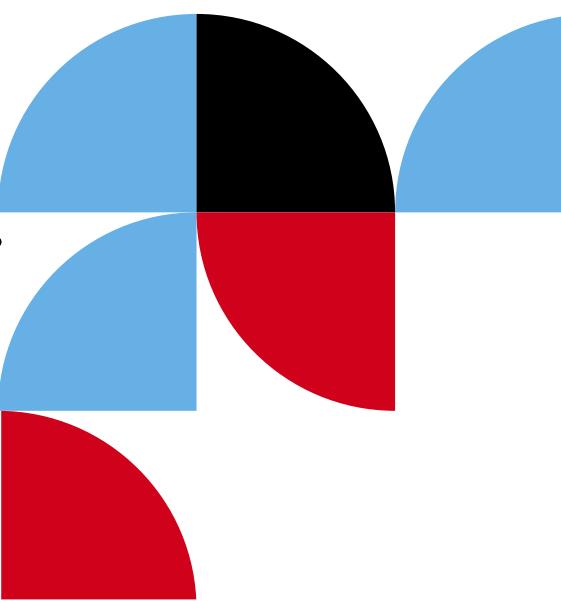
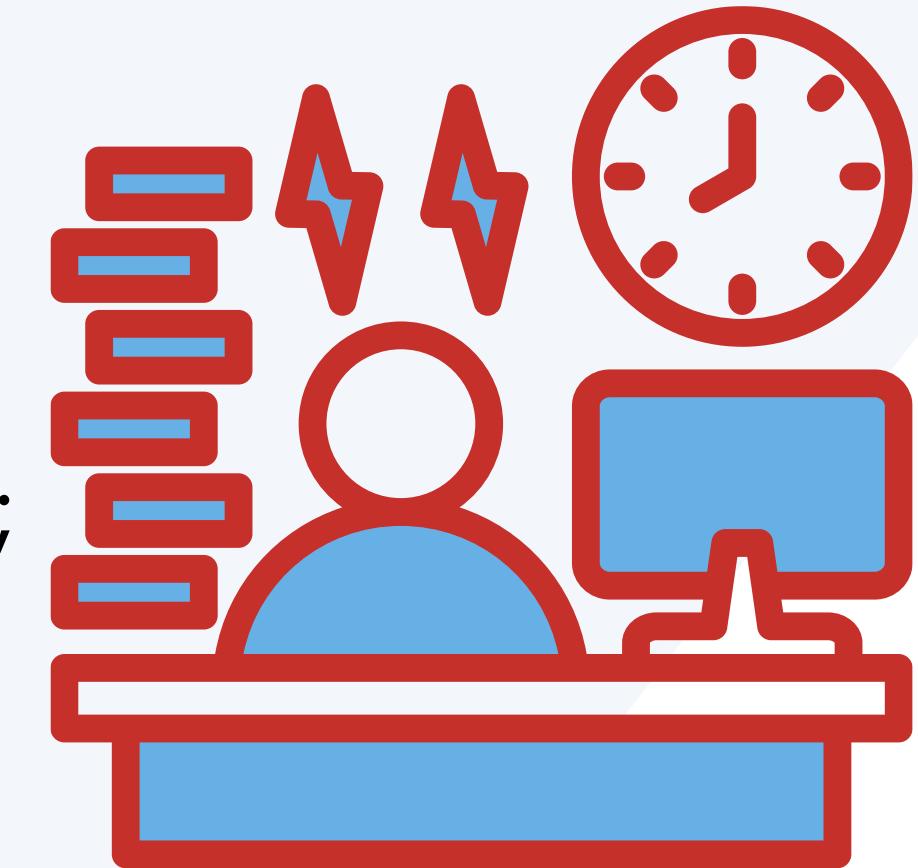
- Other parameters:
  - Number of columns generated
  - Number of threads

Workload	Proportion				
	Read	Update	Scan	Insert	Read Modify Write
A: Update heavy workload	0.5	0.5	0	0	0
B: Read mostly workload	0.95	0.05	0	0	0
C: Read only	1	0	0	0	0
D: Read latest workload	0.95	0	0	0.05	0
E: Short range	0	0	0.95	0.05	0
F: Read-modify-write	0.5	0	0	0	0.5

# Methodology

## Core workloads flow

1. Delete the data in the database (if any);
2. Load the database, using workload A's parameter file;
3. Run workload A;
4. Run workload B;
5. Run workload C;
6. Run workload F;
7. Run workload D (This workload has side-effects on the database);
8. Delete the data in the database;
9. (Re-)Load the database, using workload E's parameter file;
10. Run workload E (This workload has side-effects on the database).



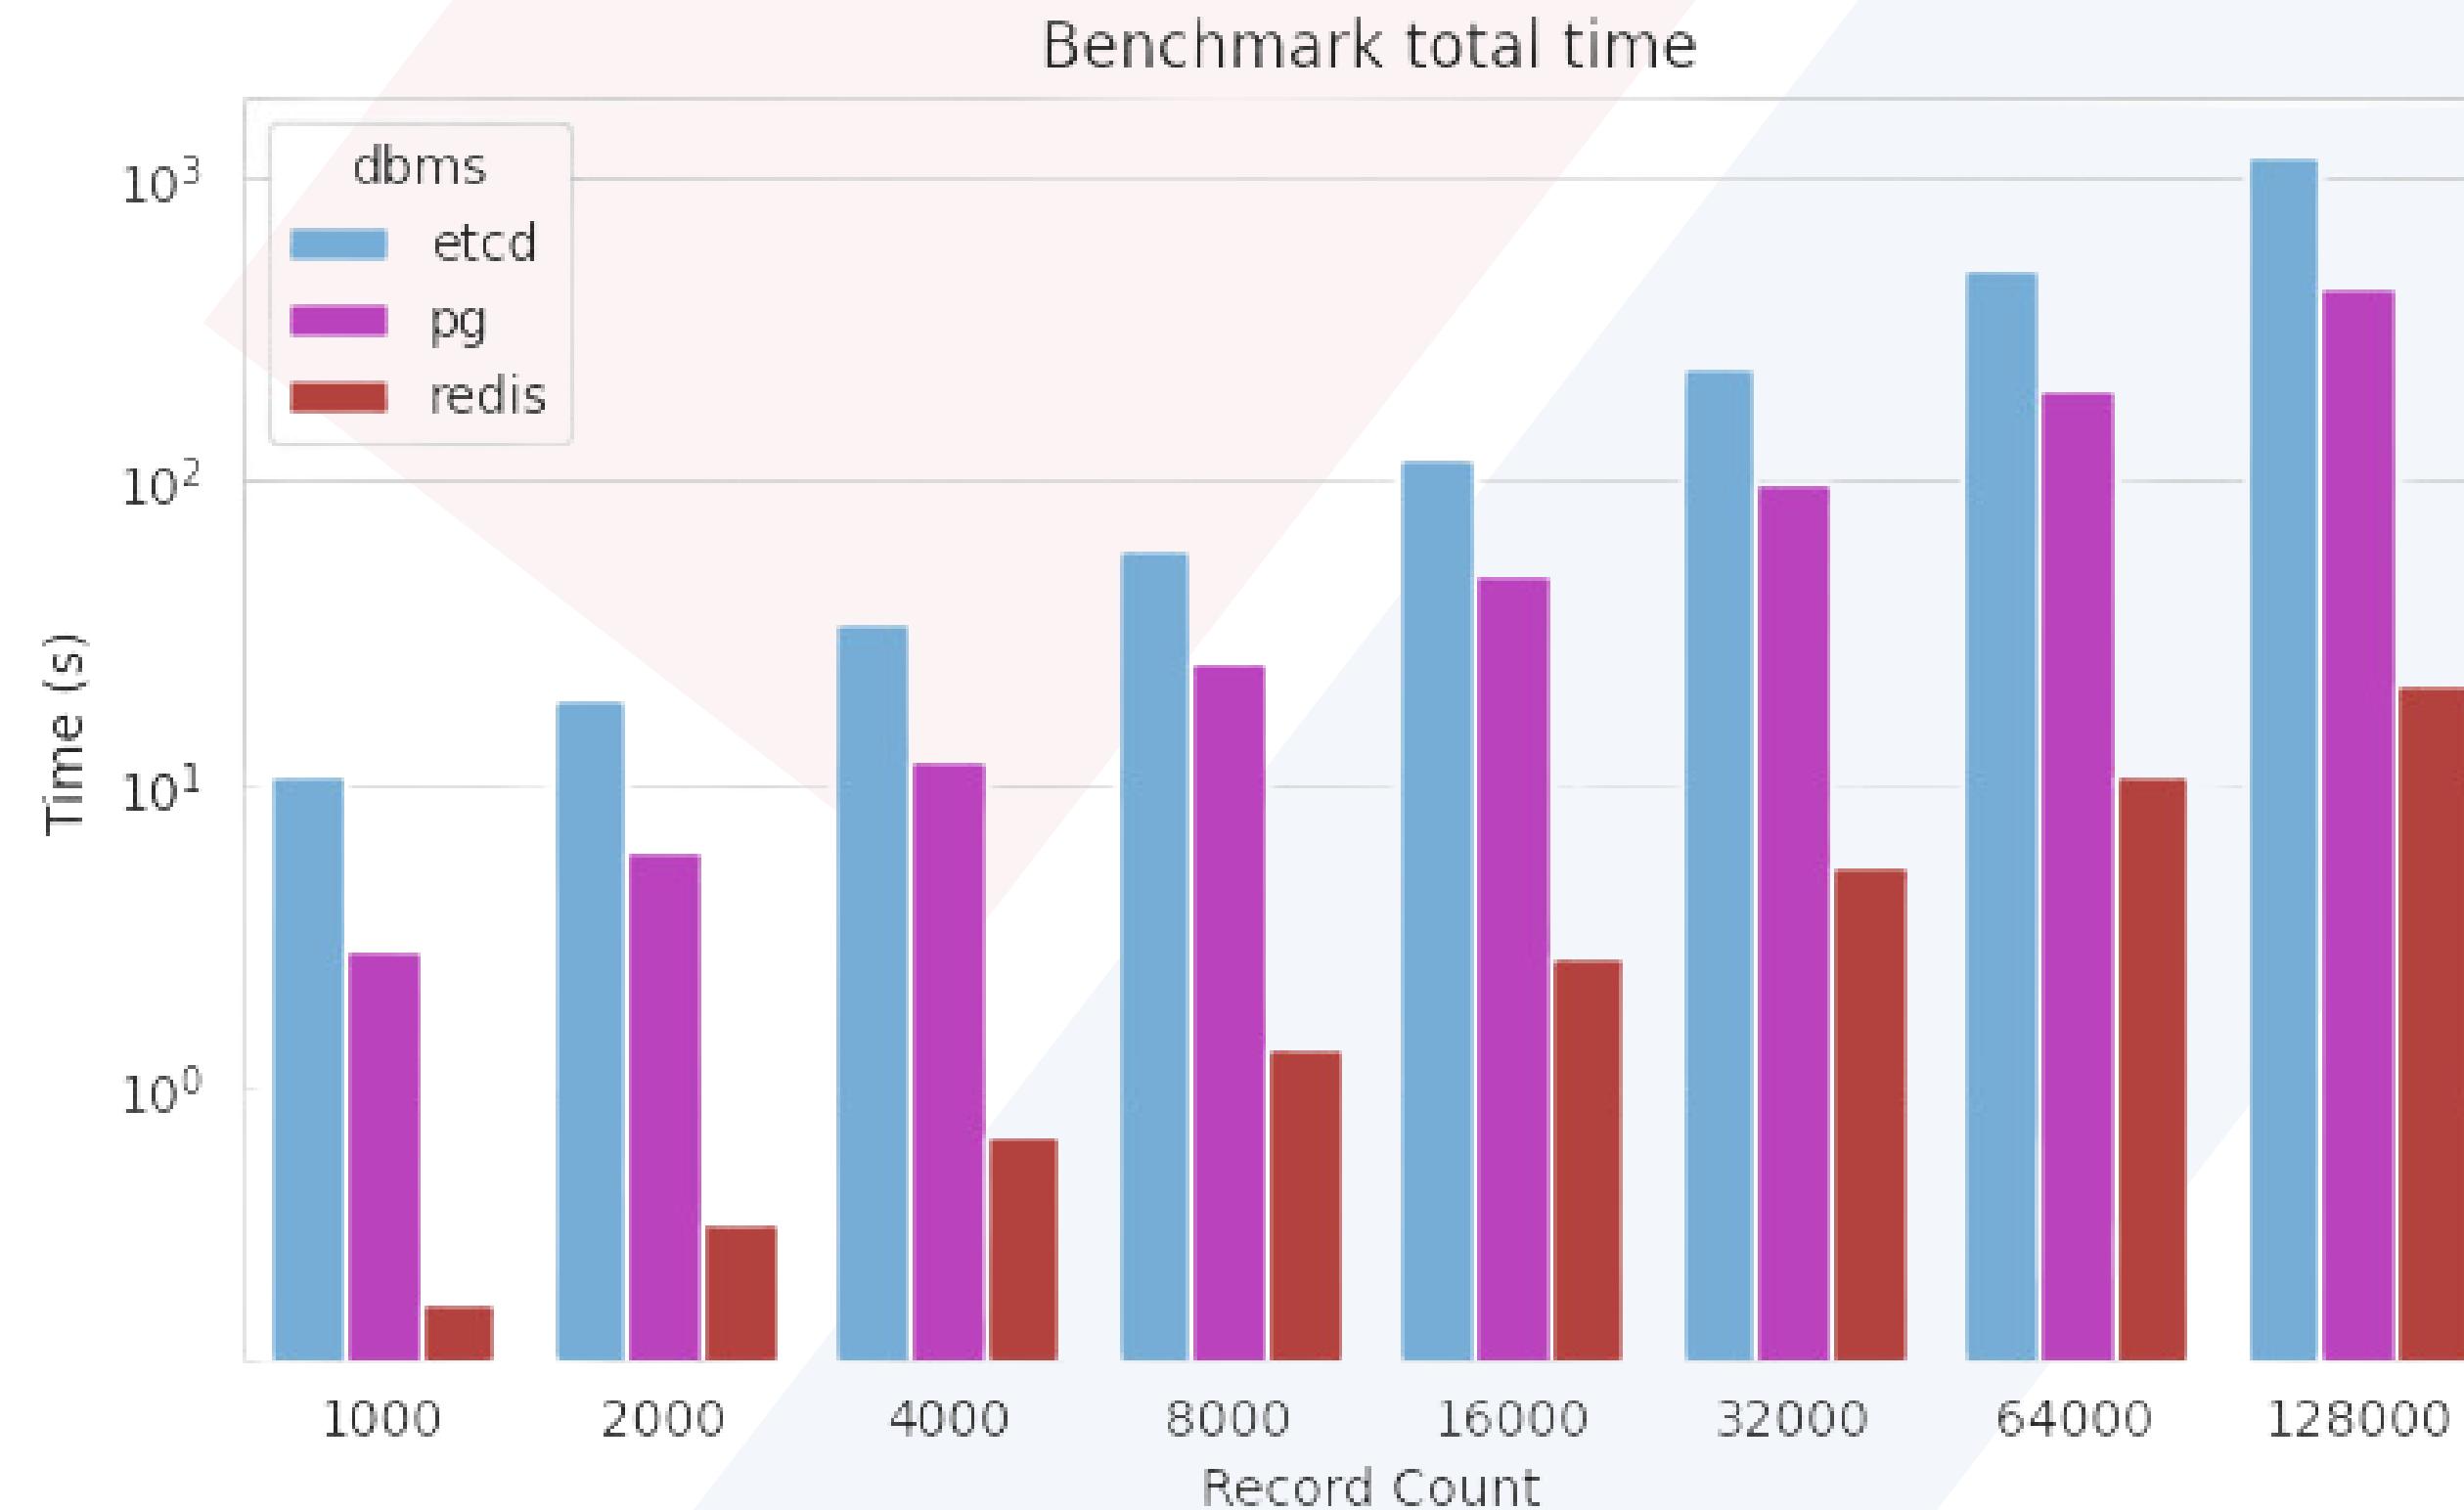
# Methodology

## Execution of Benchmark

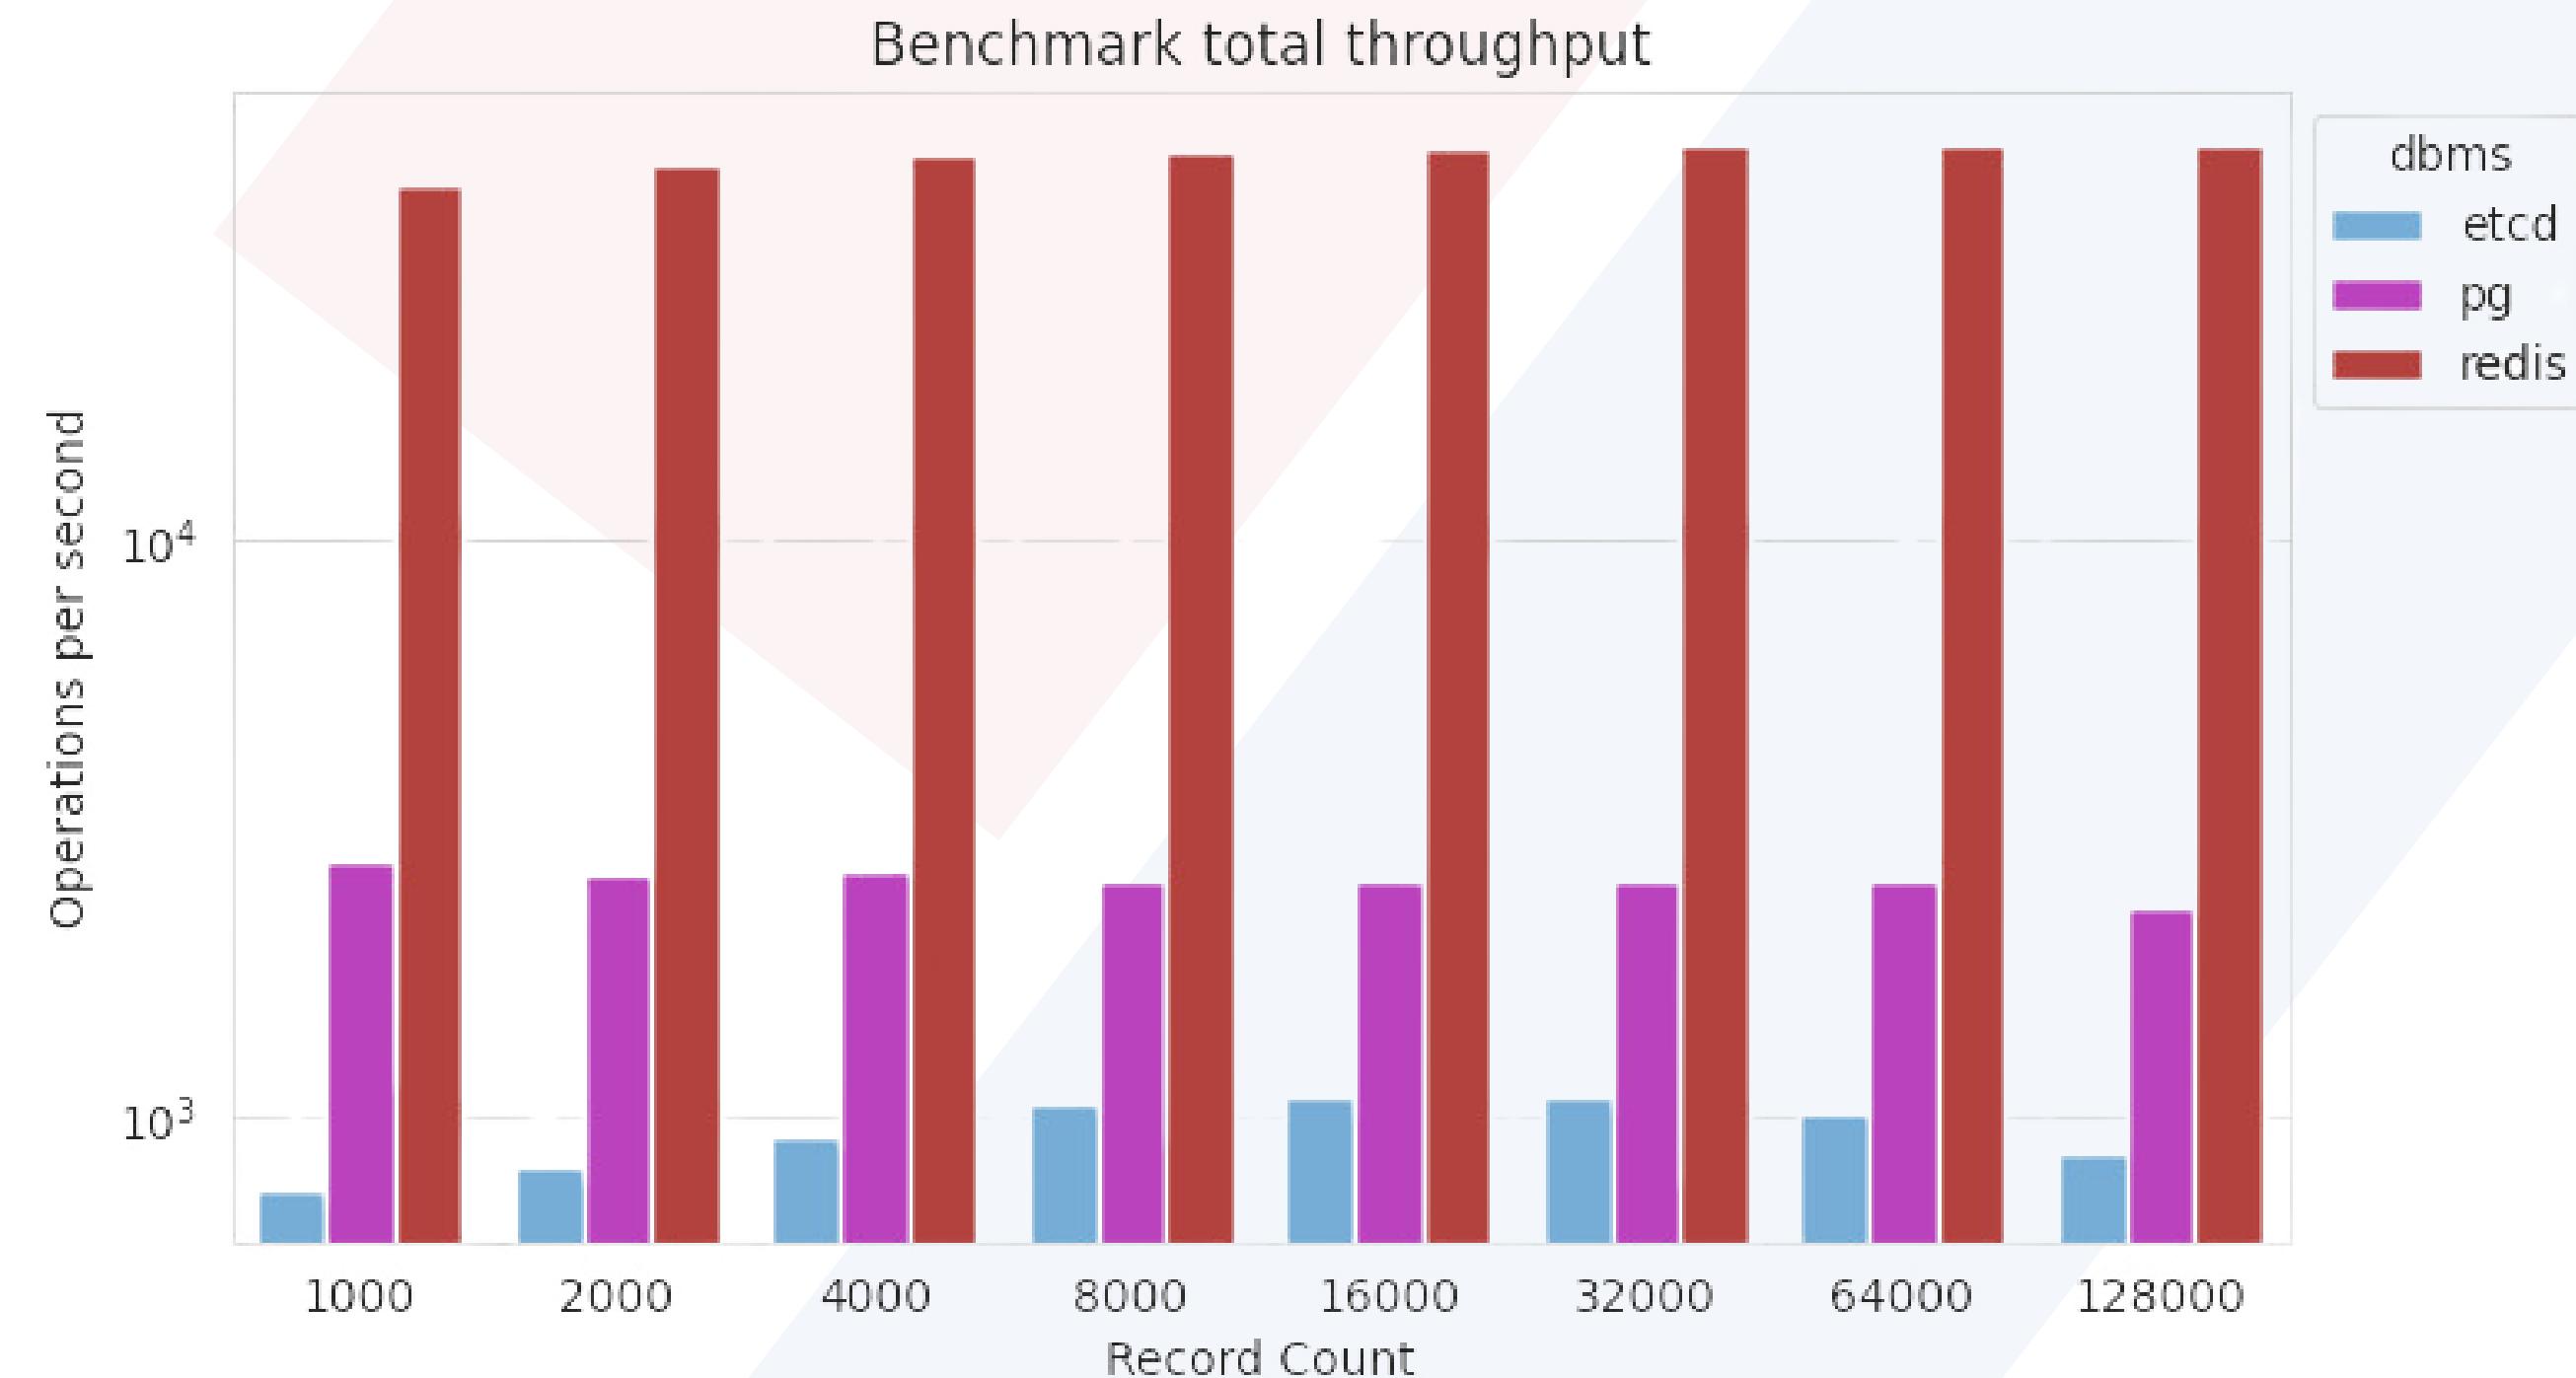
```
● ● ●

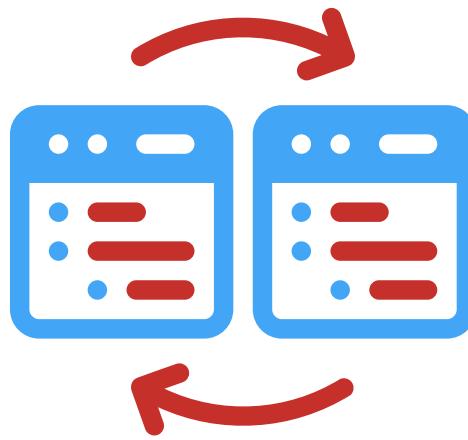
set -eu
for rep in { 1 .. 5 } ; do
    echo "Iteration ${rep}"
    for dbms in 'etcd' 'pg' 'redis' ; do
        for records in 1000 2000 4000 8000 16000 32000 64000 128000; do
            operations=${ records }
            echo -ne 'date +"%Y/%m/%d %H:%M:%S"'' ./run_benchmark ${dbms} ${operations} ${records}\n'
            ./run_benchmark.sh ${dbms} ${operations} ${records}
        done
    done
done
```

# Methodology



# Methodology



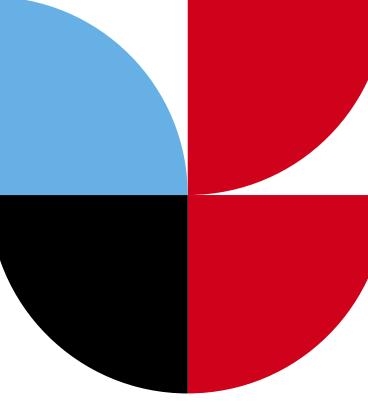


# Our Sample Console Application

*Weather & Quality of Life*

*Explorer*





# Weather & Quality of Life Explorer

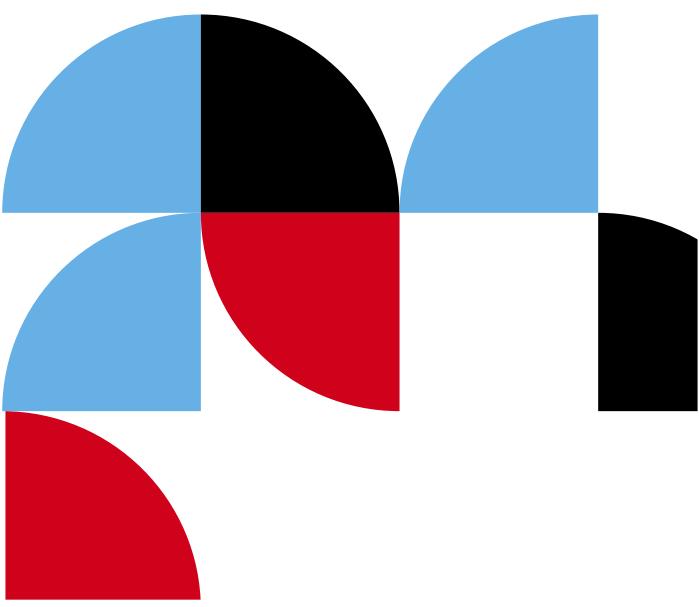
Description of application

***Sample application designed to helps users/developers to explore cities around the world.***

**Goal:** Improving execution time and performance of retrieving data when user needs to fetch data repetitively in short time.

Leverage ***caching system*** offers by key-value store DBMS

The application fetches data from [\*\*Teleport API\*\*](#).



# Weather & Quality of Life Explorer

Teleport API

OVERVIEW WIDGETS **API** SHOWCASE



## TELEPORT PUBLIC APIs

Overview

Getting Started

Resources

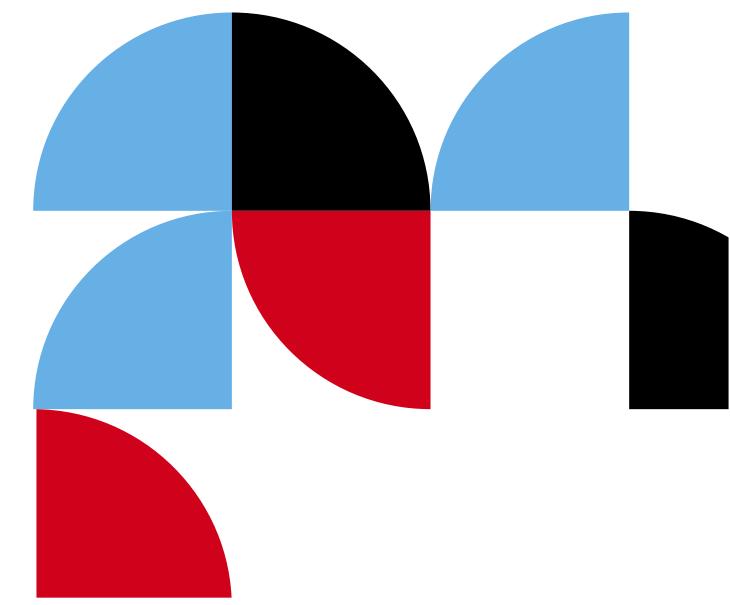
API Explorer

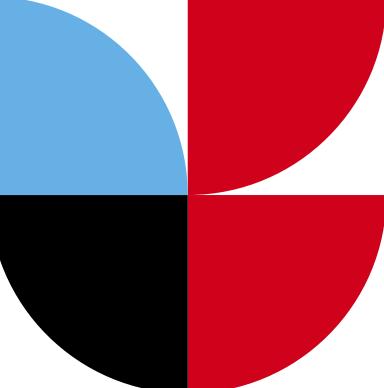
## WHY DO I NEED THIS?

If you're presenting information about various cities on your site, then you can take advantage of the Teleport API. Doesn't matter if you're running a job search site in many cities or just want to help your city stand out on its own.

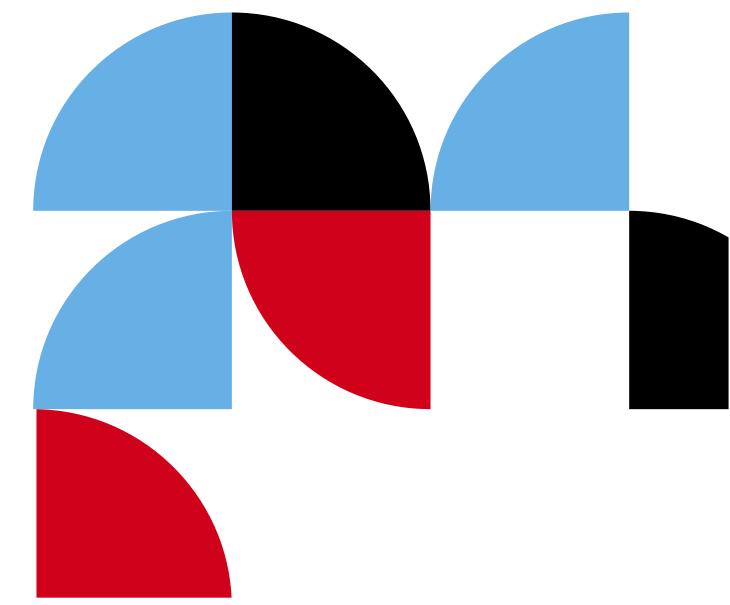
Basically anything you can see on our [Teleport Cities Pages](#) (even the photo) we can make available. If you don't see it in our API yet, ask us at [developers@teleport.org](mailto:developers@teleport.org)

- Ever needed to make sure your users provide an existing real city as an input to their location field? [City Name Search](#)
- Ever wanted to know what country or time zone a city is in? [Basic City Info](#)
- Ever wanted to present Quality of Life data (Education level, Environmental Quality, Internet Connectivity etc.) for an Urban Area? [Life Quality Data for Cities](#)
- Ever wanted to use photos of cities without the pain of building your own collection and figuring out all the rights to the images? [City Photos](#)
- Ever wanted to find cities close to user based on geo coordinates? [Nearest City Search](#)





# Weather & Quality of Life Explorer



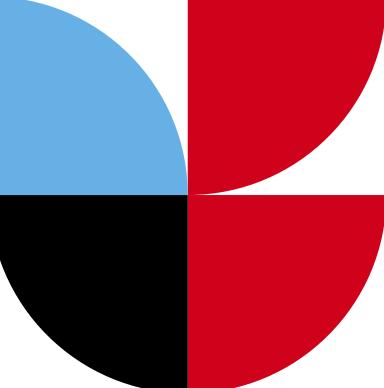
## Teleport API

### BASIC INFORMATION ABOUT A CITY

Once you have the city from the previous section, you might want to check what country or timezone it's in or how many people live there. This can easily be accomplished with the following example query for San Francisco:

Following the first linked city URL (<https://api.teleport.org/api/cities/geonameid:5391959/>) we can see details about San Francisco:

```
{  
  "links": {  
    "city:admin1_division": {  
      "href": "https://api.teleport.org/api/countries/iso_alpha2:US/admin1_divisions/geonames:CA/"  
    },  
    "city:country": {  
      "href": "https://api.teleport.org/api/countries/iso_alpha2:US/"  
    },  
    "city:timezone": {  
      "href": "https://api.teleport.org/api/timezones/iana:America%2FLos_Angeles/"  
    },  
    "city:urban_area": {  
      "href": "https://api.teleport.org/api/urban_areas/slug:san-francisco-bay-area/"  
    },  
    "self": {  
      "href": "https://api.teleport.org/api/cities/geonameid:5391959/"  
    }  
  },  
  "alternate_names": [  
    {  
      "name": "Franciscopolis"  
    },  
    {  
      "name": "Frisco"  
    },  
    ...  
  ],  
  "geoname_id": 5391959,  
  "location": {  
    "geohash": "9q8yyk8yuv26emr0cctm",  
    "latlon": {  
      "latitude": 37.77493,  
      "longitude": -122.41942  
    }  
  },  
  "name": "San Francisco",  
  "population": 805235  
}
```

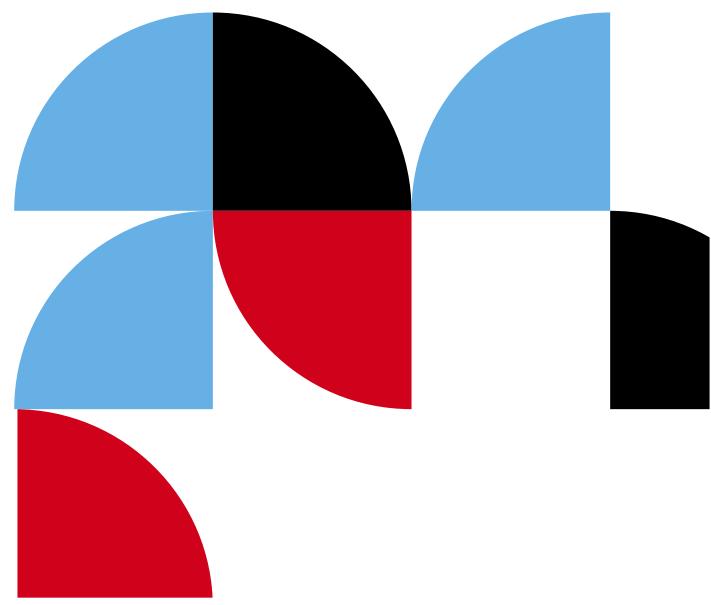


# Weather & Quality of Life Explorer

Redis integration

Fetching and displaying information *quickly and efficiently* =  
*Redis*

- *Frequently accessed city information*
- *Optimizing data retrieval*
- *Enhacing overall system efficiency*



# Weather & Quality of Life Explorer

## Redis integration

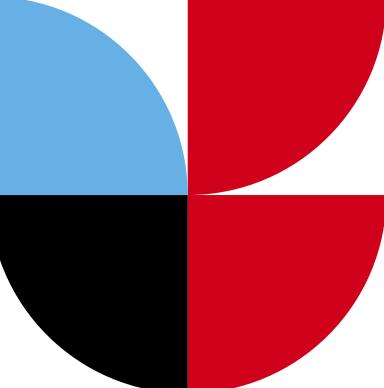
```
6  redis_client = redis.StrictRedis(host='localhost', port=6379, decode_responses=True)
7  teleport_base_url = 'https://api.teleport.org/api/'
8
9  def fetch_city_info(city_name):
10     url = f"{teleport_base_url}cities/?search={city_name}"
11     response = requests.get(url)
12     data = response.json()
13
14     if "_embedded" in data and "city:search-results" in data["_embedded"]:
15         search_results = data["_embedded"]["city:search-results"]
16         if search_results:
17             city_link = search_results[0]["_links"]["city:item"]["href"]
18             city_info_response = requests.get(city_link)
19             city_info = city_info_response.json()
20
21             # Redis caching for city information
22             ret = {
23                 'id': city_info["geoname_id"],
24                 'name': city_info["name"],
25                 'coordinates': city_info["location"]["latlon"],
26                 'population': city_info["population"]
27             }
28
29             redis_client.hset("cities", str.lower(city_name), str(ret))
30             redis_client.expire("cities", 60) # Set expiration time to 60 seconds
31             print(f"City information stored in Redis cache for {city_name}")
32
33             return ret
34
35
```

# Weather & Quality of Life Explorer

Redis integration

```
def get_cached_data():
    cached_cities = redis_client.hgetall("cities")
    cached_quality_of_life = redis_client.hgetall("quality_of_life")

    if cached_cities or cached_quality_of_life:
        return cached_cities, cached_quality_of_life
    else:
        return None, None
```

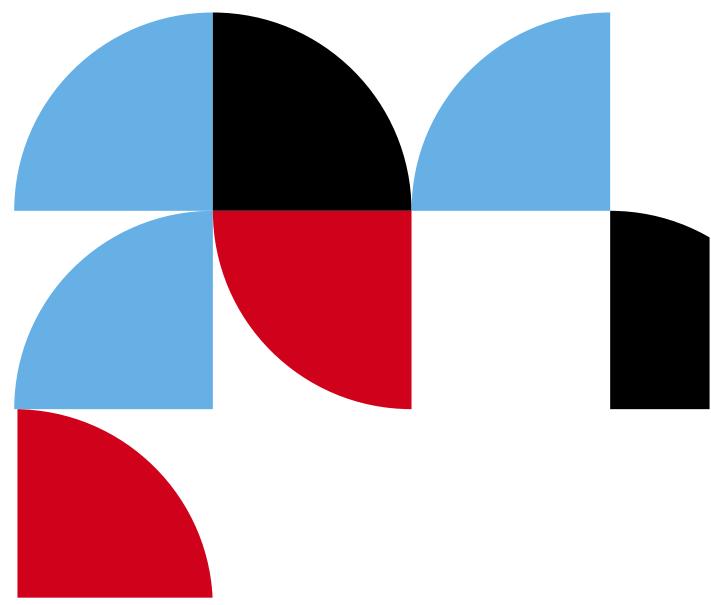


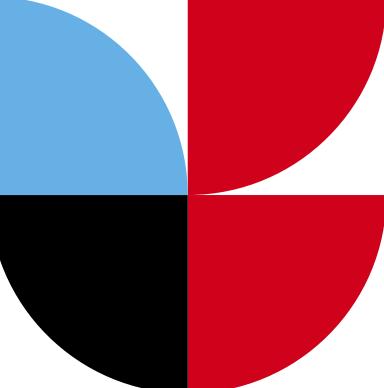
# **Weather & Quality of Life Explorer**

**Redis usage**

**Reduced API  
Calls**

**Improved responsiveness**





# Weather & Quality of Life Explorer

## Redis usage

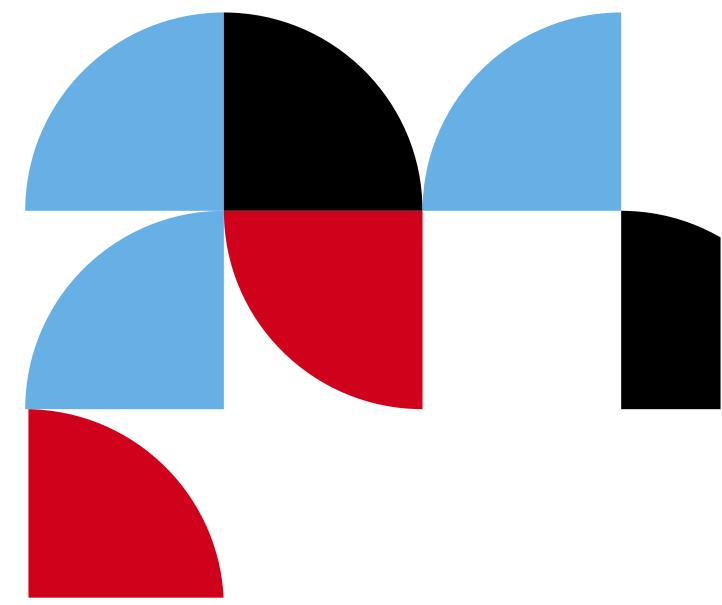
```
#####
1. Fetch City Information
2. Fetch Quality of Life Scores
3. Show Cached Cities
4. Clear Cache
5. Exit
Enter your choice (1-5): 1

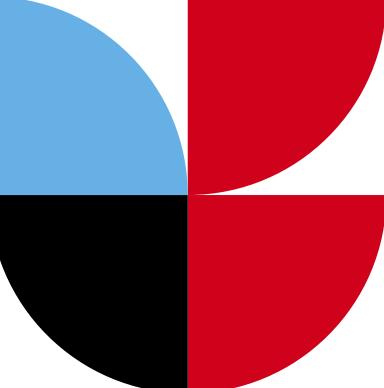
Enter the name of the city: Izmir
City information stored in Redis cache for Izmir
City information retrieved from API:
{'id': 311046, 'name': 'İzmir', 'coordinates': {'latitude': 38.41273, 'longitude': 27.13838}, 'population': 2500603}
Time taken: 0.70233 seconds

#####
1. Fetch City Information
2. Fetch Quality of Life Scores
3. Show Cached Cities
4. Clear Cache
5. Exit
Enter your choice (1-5): 1

Enter the name of the city: Izmir
City information retrieved from cache:
{'id': 311046, 'name': 'İzmir', 'coordinates': {'latitude': 38.41273, 'longitude': 27.13838}, 'population': 2500603}
Time taken: 0.00240 seconds

#####
1. Fetch City Information
2. Fetch Quality of Life Scores
3. Show Cached Cities
4. Clear Cache
5. Exit
Enter your choice (1-5): 
```



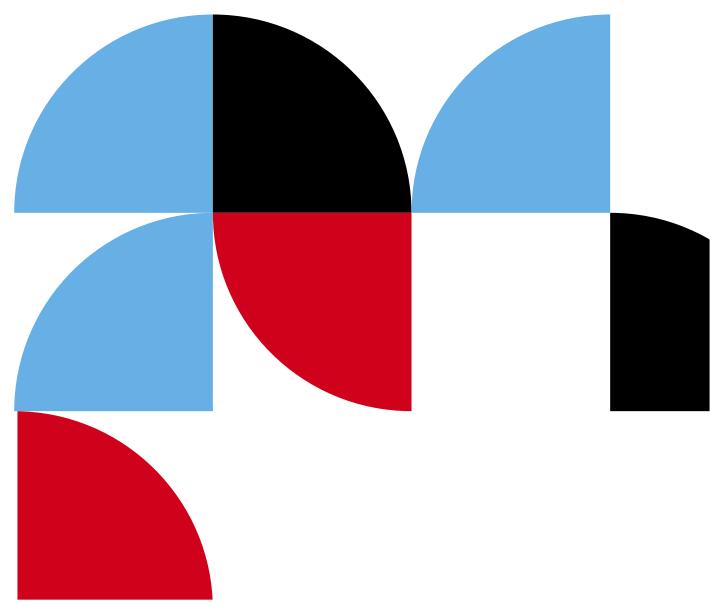


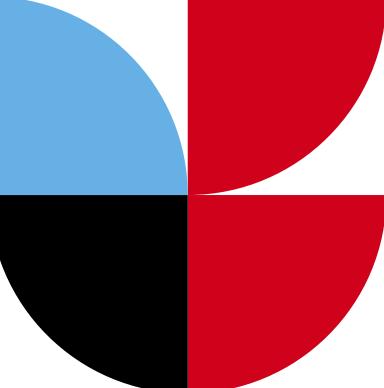
# **Weather & Quality of Life Explorer**

**Redis usage**

**Reduced API  
Calls**

**Improved responsiveness**





# Weather & Quality of Life Explorer

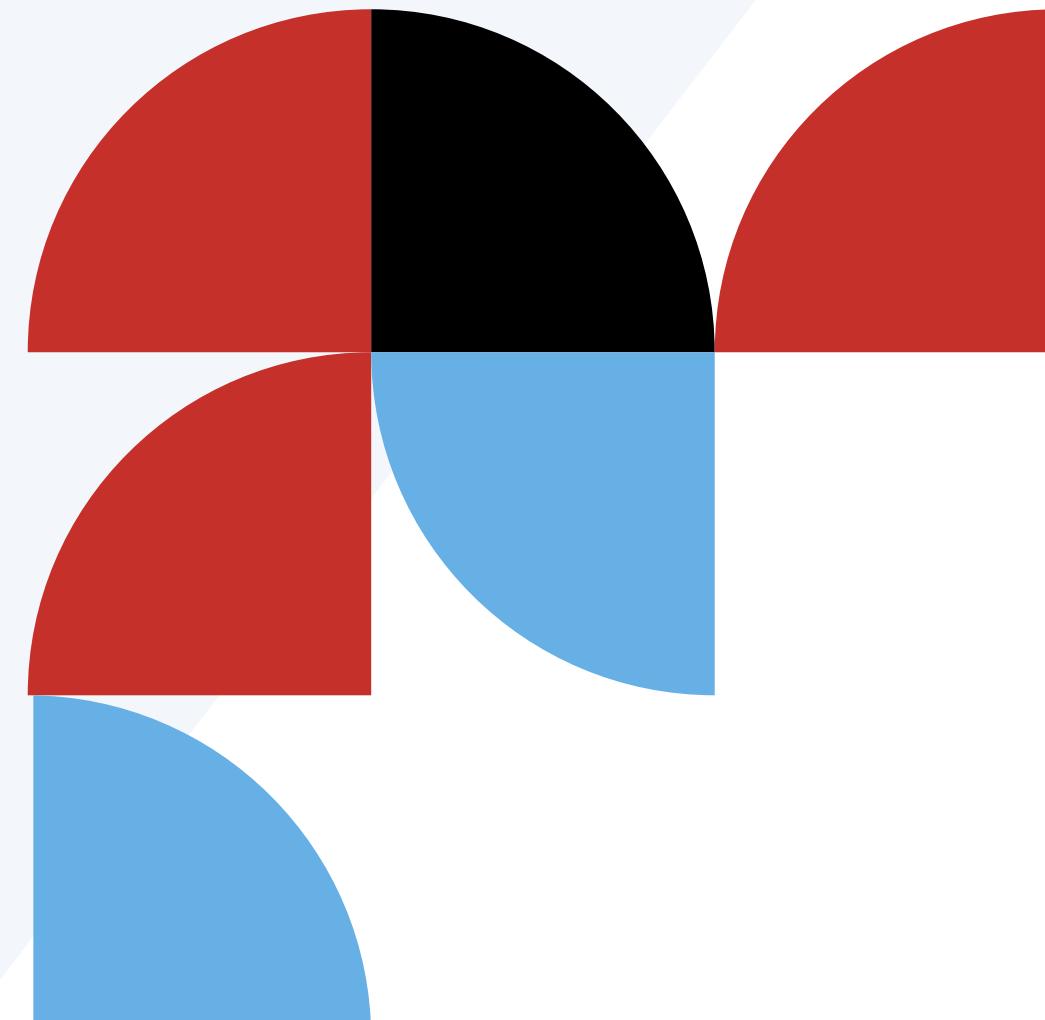
## Redis usage

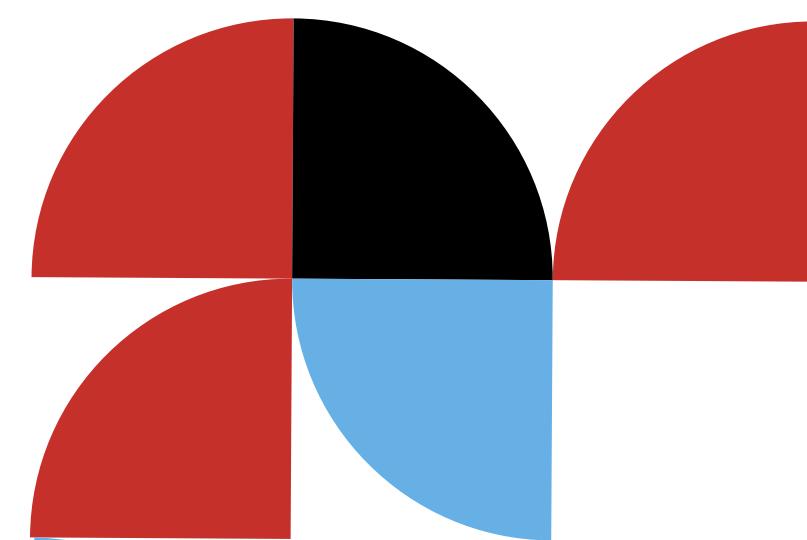
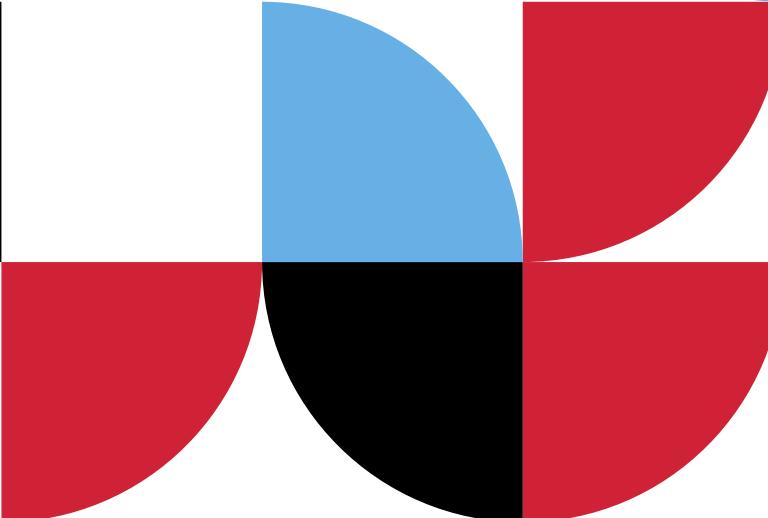
Enter the name of the urban area you want to fetch: Brussels  
Quality of Life Scores stored in Redis cache for brussels  
Quality of Life Scores:  
Housing: 6.354999999999995  
Cost of Living: 4.477  
Startups: 5.691500000000004  
Venture Capital: 3.373999999999997  
Travel Connectivity: 7.4675  
Commute: 5.386750000000001  
Business Freedom: 8.876999999999999  
Safety: 6.703  
Healthcare: 8.8163333333333  
Education: 6.653000000000005  
Environmental Quality: 4.591000000000001  
Economy: 4.811500000000006  
Taxation: 1.883000000000002  
Internet Access: 7.5715  
Leisure & Culture: 8.14399999999998  
Tolerance: 6.481  
Outdoors: 4.123  
Time taken: 0.49997 seconds ←

Enter the name of the urban area you want to fetch: Brussels  
Quality of Life Scores retrieved from cache for brussels  
Quality of Life Scores:  
Housing: 6.354999999999995  
Cost of Living: 4.477  
Startups: 5.691500000000004  
Venture Capital: 3.373999999999997  
Travel Connectivity: 7.4675  
Commute: 5.386750000000001  
Business Freedom: 8.876999999999999  
Safety: 6.703  
Healthcare: 8.8163333333333  
Education: 6.653000000000005  
Environmental Quality: 4.591000000000001  
Economy: 4.811500000000006  
Taxation: 1.883000000000002  
Internet Access: 7.5715  
Leisure & Culture: 8.14399999999998  
Tolerance: 6.481  
Outdoors: 4.123  
Time taken: 0.00093 seconds ←

# Conclusion

- Proved Redis's performance in repetitive jobs using a cache for an API.
- Redis outperformed PostgreSQL and etcd in the YCSB benchmark.
- Redis showed low-latency and efficient caching, recommended for the application.
- Future work could integrate etcd for consistency and transactional capabilities.
- The project aids developers in database selection and contributes to database management systems.





...

# **Thank You !**

+

**Let's proceed with our  
short app demo!**